

Dinis Peixoto

Phone: (+351) 966688574

GitHub: [dinispeixoto](#)

E-mail: dinispeixoto96@gmail.com

LinkedIn: [dinispeixoto](#)

Website: [dinispeixoto.com](#)

Goal-oriented Machine Learning Engineer with 4 years of professional experience. Interested in distributed systems, knowledge graphs, and machine learning engineering, in particular MLOps. Loves the challenges that come with productionizing machine learning models.

PROFESSIONAL EXPERIENCE

FARFETCH

Senior Machine Learning Engineer

Lisbon, Portugal

Sep 2018 – Present

- Design and development of the *Fashion Knowledge System*, a system that leverages an *enterprise knowledge graph* to create knowledge about business entities using ML models.
- Develop data and machine learning pipelines to load more than 500 million RDF triples into the Knowledge Graph using distributed processing tools.
- Work closely with data scientists and other engineers to productionize ML models (e.g., NLP, CV, Active Learning). Foster MLOps best practices when building automated training pipelines, model serving solutions (online and offline), orchestration, and monitoring.
- Provision infrastructure for experimentation, development, and production environments (monitoring, resources, dependency management, key vault) using Infrastructure as Code (IaC) tools and Kubernetes.
- Define company-wide guidelines for Python and Machine Learning projects. Driver of an internal community of Machine Learning Engineers.
- Conduct technical interviews for Software/Machine Learning Engineering positions. Hire, onboard, and mentor other engineers, with focus on their technical growth.
- Author of a blog post about [Mob Programming](#), internal frameworks, and an [external open-source project](#).
- **Technologies:** *Python, Spark, Databricks, MLflow, Airflow, Kubernetes, Stardog, Kafka, Cassandra, Terraform, Azure and PyTorch.*

PROJECTS

Personal blog ([dinispeixoto.com](#)): website with blog posts and quick write-ups (*Today I Learned*) about topics such as distributed systems, software, and machine learning engineering.

Brands Recommender ([GitHub](#)): system that uses product data to provide the top N recommendations for over 1000 brands available at [farfetch.com](#). *Python, TF-IDF, Cosine Similarity, GitHub Actions, and Azure ML.*

Invoice Scanner ([GitHub](#)): application that uses OCR (Optical Character Recognition) to extract data from invoice photos (e.g., products, supermarket, prices), and displays the data in a dashboard. *Tesseract OCR, ImageMagick, Ruby, ELK stack, and Docker.*

Smathome Actions - Google Assistant: Multiple Actions for Google Assistant that allow users to schedule operations and use custom invocation statements to manage smart devices. *Node.js, Python, Raspberry Pi, Actions on Google, Heroku, and GCP.*

EDUCATION

University of Minho

Master's degree, Distributed Systems and Applications Engineering

Bachelor's degree, Software Engineering

Braga, Portugal

Sep 2014 – Sep 2019

- **Vice-President at Software Engineering students union:** managed around 25 students in over 50 activities, some with hundreds of participants from different courses.
- **Mentor and Organizer at Coderdojo:** responsible for organizing monthly free programming sessions for up to 30 participants, most of them young people between 7 and 17 years old.
- **Speaker** at several talks/workshops for other students - e.g., LaTeX and Git.